

In [58]: #必要なライブラリをインポート

```
import pandas as pd
import numpy as np
np.random.seed(1337)
import pandas as pd
import keras
from keras.preprocessing.text import Tokenizer
from keras.preprocessing.sequence import pad_sequences
from keras.utils import to_categorical
from keras.layers import Dense, Input, Flatten, Dropout, Add
from keras.layers import Conv1D, MaxPooling1D, Embedding, Conv2D
from keras.models import Model
from keras.callbacks import EarlyStopping
import gensim
import nltk
from nltk.tokenize import RegexpTokenizer
from nltk.corpus import stopwords
import re
import codecs
import matplotlib.pyplot as plt
from keras.layers import Concatenate
from subprocess import check_output

stop_words = (stopwords.words('english'))[:]
stop_words.extend(["#ff", "ff", "rt"])
stop_words.remove('not')
stop_words.remove('no')
```

```
In [59]: import pandas as pd

df1 = pd.read_csv("labeled_data.csv", index_col=False)
df1.drop(['Unnamed: 0'], axis=1, inplace=True)
df1.sort_index(inplace=True)

df2 = pd.read_csv('hatespeech-train.csv', index_col=False)
df2.drop(['Unnamed: 0'], axis=1, inplace=True)
df_test = pd.read_csv('hatespeech-test.csv', index_col=False)
df_test.drop(['Unnamed: 0'], axis=1, inplace=True)
```

```
In [60]: df= pd.concat([df1.tweet,df2.tweet])
add=[]
for idx, i in enumerate(df):
    for j in df_test.tweet:
        if i==j:
            add.append(idx)
y=pd.concat([df1["class"],df2["class"]])
```

```
In [61]: if add !=[]:
    print("重複データあり")
```

重複データあり

```
In [62]: #除去
df0 = df.iloc[list(set(range(len(df.values))) - set(add))]
y= y.iloc[list(set(range(len(df.values))) - set(add))]
```

```
In [63]: df=df0
```

```
In [64]: df.reset_index(inplace=True,drop=True)
```

```
In [65]: y.reset_index(inplace=True,drop=True)
df_train= pd.concat([df, y],axis=1)
```

In []:

```
In [66]: #NaN值除去
df_train= df_train[~df_train.isnull().any(axis=1)]
y= y[~df_train.isnull().any(axis=1)]
```

```
In [67]: df_train.columns=['text','class']
```

```
In [68]: df_test.columns= ["count","hate_speech","offensive_language","neither","class","text"]
```

```
In [69]: df_train[df_train["class"]==2].count()
```


```
Out[69]: text      7375
class      7375
dtype: int64
```

```
In [70]: add0= df_train[df_train["class"]==0].sample(20646,replace=True)
add2= df_train[df_train["class"]==2].sample(15825,replace=True)
```

```
In [71]: df_train= pd.concat([df_train,add0,add2])
```

```
In [72]: df_train["class"].hist()
```

```
Out[72]: <matplotlib.axes._subplots.AxesSubplot at 0x1b69b44080>
```



```
In [73]: from keras.utils import to_categorical
encoded = to_categorical(df_train['class'])
encoded2 = to_categorical(df_test['class'])
y_e = pd.DataFrame(encoded)
y_e.columns=["hate_speech", "offensive_language", "neither"]
y_e2 = pd.DataFrame(encoded2)
y_e2.columns=["hate_speech", "offensive_language", "neither"]
df_train.reset_index(inplace=True, drop=True)
df_train = pd.concat([df_train["text"], y_e], axis=1)
df_test = pd.concat([df_test["text"], y_e2], axis=1)
```

```
In [ ]:
```

```
In [ ]:
```

```
In [74]: EMBEDDING_DIM = 300
MAX_VOCAB_SIZE = 24000
MAX_SEQUENCE_LENGTH = 83

batch_size = 64
```

```
In [75]: def standardize_text(df, text_field):
df[text_field] = df[text_field].str.replace(r"http\S+", "")
df[text_field] = df[text_field].str.replace(r"http", "")
df[text_field] = df[text_field].str.replace(r"\S+", "")
df[text_field] = df[text_field].str.replace(r"^[^A-Za-z0-9(),!@'\\". _\n]", " ")
df[text_field] = df[text_field].str.replace(r"@", "")
df[text_field] = df[text_field].str.replace(r"!", "")
df[text_field] = df[text_field].str.lower()
df[text_field] = df[text_field].str.replace(r"http[s]?://(?:[a-zA-Z]|[0-9]|[$-_@.&+]|'[*\(\),]|(?:%[0-9a-fA-F]{2})")
df[text_field] = df[text_field].str.replace(r"@[\w-]+", " ")
df[text_field] = df[text_field].str.replace(r'\s+', ' ')
df[text_field] = df[text_field].str.replace(r'^[! ]*RT', '')
df[text_field] = df[text_field].str.replace(r'^[! ]*rt', '')
return df
```

```
In [76]: train_comments = standardize_text(df_train, "text")
train_comments.to_csv("train_clean_data.csv")
train_comments.head()
```

```
Out[76]:
```

	text	hate_speech	offensive_language	neither
0	as a woman you shouldn't complain about clean...	0.0	0.0	1.0
1	boy dats cold tyga dwn bad for cuffin dat hoe...	0.0	1.0	0.0
2	dawg rt you ever fuck a bitch and she start t...	0.0	1.0	0.0
3	she look like a tranny	0.0	1.0	0.0
4	the shit you hear about me might be true or i...	0.0	1.0	0.0

```
In [ ]:
```

```
In [77]: label_names = ["hate_speech", "offensive_language", "neither"]
y_train = train_comments[label_names].values
```

```
In [78]: test_comments = standardize_text(df_test, "text")
test_comments.to_csv("test_clean_data.csv")
test_comments.head()
```

```
Out[78]:
```

	text	hate_speech	offensive_language	neither
0	still really confused re humplayproblems i thi...	1.0	0.0	0.0
1	stoni is a fuckin queer	1.0	0.0	0.0
2	stu a rape faced faggot man i'm sick of this n...	1.0	0.0	0.0
3	stuck up serious bitches are possibly the wors...	1.0	0.0	0.0
4	stupid teabagger restaurant making customers p...	1.0	0.0	0.0

```
In [79]: label_names = ["hate_speech", "offensive_language", "neither"]
y_test = test_comments[label_names].values
```

```
In [ ]:
```

```
In [80]: tokenizer = RegexpTokenizer(r'\w+')
clean_train_comments = pd.read_csv("train_clean_data.csv")
clean_train_comments['text'] = clean_train_comments['text'].astype('str')
clean_train_comments.dtypes
clean_train_comments["tokens"] = clean_train_comments["text"].apply(tokenizer.tokenize)

clean_train_comments["tokens"] = clean_train_comments["tokens"].apply(lambda vec: [word for word in vec
clean_train_comments.head()
```

Out[80]:

	Unnamed: 0	text	hate_speech	offensive_language	neither	tokens
0	0	as a woman you shouldn't complain about clean...	0.0	0.0	1.0	[woman, complain, cleaning, house, amp, man, a...
1	1	boy dats cold tyga dwn bad for cuffin dat hoe...	0.0	1.0	0.0	[boy, dats, cold, tyga, dwn, bad, cuffin, dat,...
2	2	dawg rt you ever fuck a bitch and she start t...	0.0	1.0	0.0	[dawg, ever, fuck, bitch, start, cry, confused...
3	3	she look like a tranny	0.0	1.0	0.0	[look, like, tranny]
4	4	the shit you hear about me might be true or i...	0.0	1.0	0.0	[shit, hear, might, true, might, faker, bitch,...

```
In [81]: clean_test_comments = pd.read_csv("test_clean_data.csv")
clean_test_comments['text'] = clean_test_comments['text'].astype('str')
clean_test_comments.dtypes
clean_test_comments["tokens"] = clean_test_comments["text"].apply(tokenizer.tokenize)
clean_test_comments["tokens"] = clean_test_comments["tokens"].apply(lambda vec: [word for word in vec if
clean_test_comments.head()
```

Out[81]:

	Unnamed: 0	text	hate_speech	offensive_language	neither	tokens
0	0	still really confused re humplayproblems i thi...	1.0	0.0	0.0	[still, really, confused, humplayproblems, thi...
1	1	stoni is a fuckin queer	1.0	0.0	0.0	[stoni, fuckin, queer]
2	2	stu a rape faced faggot man i'm sick of this n...	1.0	0.0	0.0	[stu, rape, faced, faggot, man, sick, nigga]
3	3	stuck up serious bitches are possibly the wors...	1.0	0.0	0.0	[stuck, serious, bitches, possibly, worst, typ...
4	4	stupid teabagger restaurant making customers p...	1.0	0.0	0.0	[stupid, teabagger, restaurant, making, custom...

```
In [82]: all_training_words = [word for tokens in clean_train_comments["tokens"] for word in tokens]
training_sentence_lengths = [len(tokens) for tokens in clean_train_comments["tokens"]]
TRAINING_VOCAB = sorted(list(set(all_training_words)))
print("%s 単語数, 語彙総数 %s" % (len(all_training_words), len(TRAINING_VOCAB)))
print("文の最大長さ %s" % max(training_sentence_lengths))
```

581074 単語数, 語彙総数 20431
文の最大長さ 83


```
In [83]: all_test_words = [word for tokens in clean_test_comments["tokens"] for word in tokens]
test_sentence_lengths = [len(tokens) for tokens in clean_test_comments["tokens"]]
TEST_VOCAB = sorted(list(set(all_test_words)))
print("%s 単語数, 語彙総数 %s" % (len(all_test_words), len(TEST_VOCAB)))
print("文の最大長さ %s" % max(test_sentence_lengths))
```

7699 単語数, 語彙総数 2962
文の最大長さ 23

```
In [84]: word2vec_path = "GoogleNews-vectors-negative300.bin"
word2vec = gensim.models.KeyedVectors.load_word2vec_format(word2vec_path, binary=True)

def get_average_word2vec(tokens_list, vector, generate_missing=False, k=300):
    if len(tokens_list)<1:
        return np.zeros(k)
    if generate_missing:
        vectorized = [vector[word] if word in vector else np.random.rand(k) for word in tokens_list]
    else:
        vectorized = [vector[word] if word in vector else np.zeros(k) for word in tokens_list]
    length = len(vectorized)
    summed = np.sum(vectorized, axis=0)
    averaged = np.divide(summed, length)
    return averaged
```

In []:

```
In [85]: tokenizer = Tokenizer(num_words=MAX_VOCAB_SIZE, lower=True, char_level=False)
tokenizer.fit_on_texts(clean_train_comments["text"].tolist())
training_sequences = tokenizer.texts_to_sequences(clean_train_comments["text"].tolist())

train_word_index = tokenizer.word_index
print('%s ユニックなトークン' % len(train_word_index))

train_cnn_data = pad_sequences(training_sequences, maxlen=MAX_SEQUENCE_LENGTH)

train_embedding_weights = np.zeros((len(train_word_index)+1, EMBEDDING_DIM))
for word, index in train_word_index.items():
    train_embedding_weights[index,:] = word2vec[word] if word in word2vec else np.random.rand(EMBEDDING_DIM)
print(train_embedding_weights.shape)

21205 ユニックなトークン
(21206, 300)
```

```
In [86]: test_sequences = tokenizer.texts_to_sequences(clean_test_comments["text"].tolist())
test_cnn_data = pad_sequences(test_sequences, maxlen=MAX_SEQUENCE_LENGTH)
```

```

In [87]: def ConvNet(embeddings, max_sequence_length, num_words, embedding_dim, labels_index, trainable=False):

    embedding_layer = Embedding(num_words,
                                embedding_dim,
                                weights=[embeddings],
                                input_length=max_sequence_length,
                                trainable=trainable)

    sequence_input = Input(shape=(max_sequence_length,), dtype='int32')
    embedded_sequences = embedding_layer(sequence_input)

    convs = []
    filter_sizes = [3,4,5]

    for filter_size in filter_sizes:
        l_conv = Conv1D(filters=128, kernel_size=filter_size, activation='relu')(embedded_sequences)
        l_pool = MaxPooling1D(pool_size=3)(l_conv)
        convs.append(l_pool)

    l_merge = Concatenate(axis=1)(convs)
    x = Flatten()(l_merge)
    x = Dense(128, activation='relu')(x)
    preds = Dense(labels_index, activation='softmax')(x)

    model = keras.models.Model(sequence_input, preds)
    model.compile(loss='categorical_crossentropy',
                  optimizer='rmsprop',
                  metrics=['acc'])
    model.summary()
    return model

```

```

In [88]: x_train = train_cnn_data
         x_test = test_cnn_data

```

```
In [89]: model = ConvNet(train_embedding_weights, MAX_SEQUENCE_LENGTH, len(train_word_index)+1, EMBEDDING_DIM,
                        len(list(label_names)), False)
```

Layer (type)	Output Shape	Param #	Connected to
input_5 (InputLayer)	(None, 83)	0	
embedding_5 (Embedding)	(None, 83, 300)	6361800	input_5[0][0]
conv1d_10 (Conv1D)	(None, 81, 128)	115328	embedding_5[0][0]
conv1d_11 (Conv1D)	(None, 80, 128)	153728	embedding_5[0][0]
conv1d_12 (Conv1D)	(None, 79, 128)	192128	embedding_5[0][0]
max_pooling1d_10 (MaxPooling1D)	(None, 27, 128)	0	conv1d_10[0][0]
max_pooling1d_11 (MaxPooling1D)	(None, 26, 128)	0	conv1d_11[0][0]
max_pooling1d_12 (MaxPooling1D)	(None, 26, 128)	0	conv1d_12[0][0]
concatenate_4 (Concatenate)	(None, 79, 128)	0	max_pooling1d_10[0][0] max_pooling1d_11[0][0] max_pooling1d_12[0][0]
flatten_3 (Flatten)	(None, 10112)	0	concatenate_4[0][0]
dense_7 (Dense)	(None, 128)	1294464	flatten_3[0][0]
dense_8 (Dense)	(None, 3)	387	dense_7[0][0]
Total params: 8,117,835			
Trainable params: 1,756,035			
Non-trainable params: 6,361,800			

```
In [90]: hist = model.fit(x_train, y_train, epochs=10, validation_split=0.1, shuffle=True, batch_size=batch_size)
```

Train on 62640 samples, validate on 6960 samples

Epoch 1/10

62640/62640 [=====] - 150s 2ms/step - loss: 0.3029 - acc: 0.8826 - val_loss: 0.1782 - val_acc: 0.9412

Epoch 2/10

62640/62640 [=====] - 146s 2ms/step - loss: 0.0796 - acc: 0.9749 - val_loss: 0.0239 - val_acc: 0.9938

Epoch 3/10

62640/62640 [=====] - 164s 3ms/step - loss: 0.0489 - acc: 0.9865 - val_loss: 0.0138 - val_acc: 0.9970

Epoch 4/10

62640/62640 [=====] - 145s 2ms/step - loss: 0.0365 - acc: 0.9898 - val_loss: 0.0094 - val_acc: 0.9980

Epoch 5/10

62640/62640 [=====] - 146s 2ms/step - loss: 0.0294 - acc: 0.9917 - val_loss: 0.0079 - val_acc: 0.9980

Epoch 6/10

62640/62640 [=====] - 152s 2ms/step - loss: 0.0265 - acc: 0.9933 - val_loss: 0.0342 - val_acc: 0.9879

Epoch 7/10

62640/62640 [=====] - 172s 3ms/step - loss: 0.0212 - acc: 0.9945 - val_loss: 0.0198 - val_acc: 0.9953

Epoch 8/10

62640/62640 [=====] - 137s 2ms/step - loss: 0.0224 - acc: 0.9945 - val_loss: 0.0188 - val_acc: 0.9955

Epoch 9/10

62640/62640 [=====] - 136s 2ms/step - loss: 0.0180 - acc: 0.9953 - val_loss: 3.7572e-04 - val_acc: 1.0000

Epoch 10/10

62640/62640 [=====] - 142s 2ms/step - loss: 0.0180 - acc: 0.9956 - val_loss: 0.0204 - val_acc: 0.9948

```
In [91]: y_pred = model.predict(x_test, batch_size=1024, verbose=1)
```

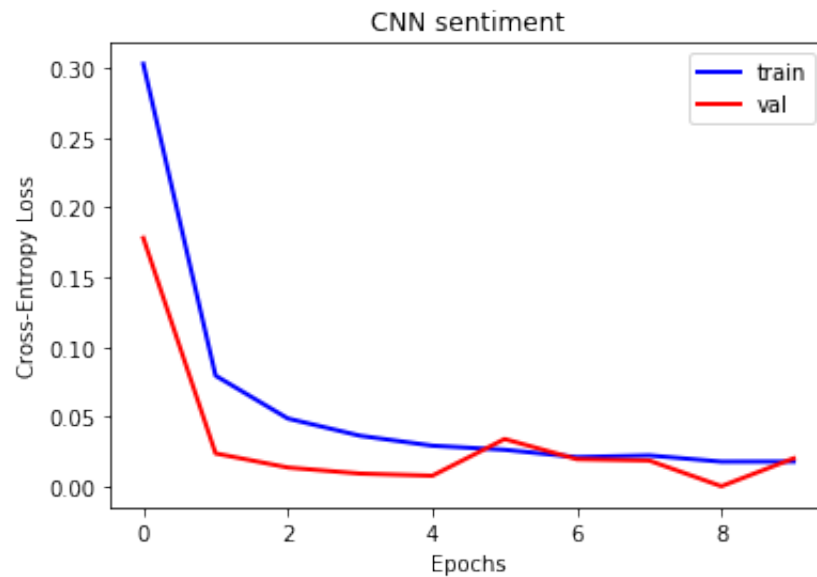
1040/1040 [=====] - 2s 2ms/step

In [92]:


```
submission_df = pd.DataFrame(columns=['id'] + label_names)
submission_df['id'] = test_comments.index.values
submission_df[label_names] = y_pred
submission_df.to_csv("./cnn_submission.csv", index=False)
```

In [93]:

```
plt.figure()
plt.plot(hist.history['loss'], lw=2.0, color='b', label='train')
plt.plot(hist.history['val_loss'], lw=2.0, color='r', label='val')
plt.title('CNN sentiment')
plt.xlabel('Epochs')
plt.ylabel('Cross-Entropy Loss')
plt.legend(loc='upper right')
plt.show()
```



```
In [94]: plt.figure()
plt.plot(hist.history['acc'], lw=2.0, color='b', label='train')
plt.plot(hist.history['val_acc'], lw=2.0, color='r', label='val')
plt.title('CNN sentiment')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend(loc='upper left')
plt.show()
```




```
In [95]: from sklearn.metrics import classification_report
report = classification_report(np.argmax(y_test,axis =1), np.argmax(y_pred,axis =1))
```

```
In [96]: print(report)
```

	precision	recall	f1-score	support
0	0.46	0.35	0.40	140
1	0.73	0.90	0.81	500
2	0.95	0.76	0.84	400
micro avg	0.77	0.77	0.77	1040
macro avg	0.71	0.67	0.68	1040
weighted avg	0.78	0.77	0.76	1040


```
In [97]: import seaborn
from sklearn.metrics import confusion_matrix
confusion_matrix = confusion_matrix(np.argmax(y_test,axis =1),np.argmax(y_pred,axis =1))
matrix_proportions = np.zeros((3,3))
for i in range(0,3):
    matrix_proportions[i,:] = confusion_matrix[i,:]/float(confusion_matrix[i,:].sum())
names=['Hate','Offensive','Neither']
confusion_df = pd.DataFrame(matrix_proportions, index=names,columns=names)
plt.figure(figsize=(5, 5))
seaborn.heatmap(confusion_df,annot=True,annot_kws={"size": 12},cmap='YlGnBu',cbar=False, square=True,fmt
plt.ylabel(r'True categories',fontsize=14)
plt.xlabel(r'Predicted categories',fontsize=14)
plt.tick_params(labelsize=12)
```



In []:

In []:


```

        bias=True, **kwargs):
    self.supports_masking = True
    self.init = initializers.get('glorot_uniform')

    self.W_regularizer = regularizers.get(W_regularizer)
    self.b_regularizer = regularizers.get(b_regularizer)

    self.W_constraint = constraints.get(W_constraint)
    self.b_constraint = constraints.get(b_constraint)

    self.bias = bias
    self.step_dim = step_dim
    self.features_dim = 0
    super(Attention, self).__init__(**kwargs)

def build(self, input_shape):
    assert len(input_shape) == 3

    self.W = self.add_weight((input_shape[-1],),
                              initializer=self.init,
                              name='{}_W'.format(self.name),
                              regularizer=self.W_regularizer,
                              constraint=self.W_constraint)
    self.features_dim = input_shape[-1]

    if self.bias:
        self.b = self.add_weight((input_shape[1],),
                                  initializer='zero',
                                  name='{}_b'.format(self.name),
                                  regularizer=self.b_regularizer,
                                  constraint=self.b_constraint)
    else:
        self.b = None

    self.built = True

def compute_mask(self, input, input_mask=None):
    return None

def call(self, x, mask=None):

```

```

def call(self, x, mask=None):
    features_dim = self.features_dim
    step_dim = self.step_dim

    eij = K.reshape(K.dot(K.reshape(x, (-1, features_dim)),
                          K.reshape(self.W, (features_dim, 1))), (-1, step_dim))

    if self.bias:
        eij += self.b

    eij = K.tanh(eij)

    a = K.exp(eij)

    if mask is not None:
        a *= K.cast(mask, K.floatx())

    a /= K.cast(K.sum(a, axis=1, keepdims=True) + K.epsilon(), K.floatx())

    a = K.expand_dims(a)
    weighted_input = x * a
    return K.sum(weighted_input, axis=1)

def compute_output_shape(self, input_shape):
    return input_shape[0], self.features_dim

```

```

def LSTM(embeddings, max_sequence_length, num_words, embedding_dim, labels_index, trainable=False):

```

```

    #パラメータ
    act = 'relu'
    re_weight = True
    num_lstm = 300
    num_dense = 150

    sequence_1_input = Input(shape=(max_sequence_length,), dtype='int32')

    embedding_layer = Embedding(num_words,
                                embedding_dim,

```

```

        weights=[embeddings],
        input_length=max_sequence_length,
        trainable=trainable)

embedded_sequences_1 = embedding_layer(sequence_1_input)

lstm_layer = LSTM(num_lstm, return_sequences=True)

x1 = lstm_layer(embedded_sequences_1)

merged = Attention(max_sequence_length)(x1)

merged = Dense(num_dense, activation=act)(merged)

preds = Dense(3, activation='softmax')(merged)


model = Model(sequence_1_input, preds)
model.compile(loss='categorical_crossentropy',
              optimizer='adam',
              metrics=['acc'])
model.summary()
return model

```

```
In [35]: model = Lstm(train_embedding_weights, MAX_SEQUENCE_LENGTH, len(train_word_index)+1, EMBEDDING_DIM,
                    len(list(label_names)), False)
```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	(None, 83)	0
embedding_1 (Embedding)	(None, 83, 300)	6361800
lstm_1 (LSTM)	(None, 83, 300)	721200
attention_1 (Attention)	(None, 300)	383
dense_1 (Dense)	(None, 150)	45150
dense_2 (Dense)	(None, 3)	453
Total params: 7,128,986		
Trainable params: 767,186		
Non-trainable params: 6,361,800		

```
In [94]: hist = model.fit(x_train, y_train, epochs=10, validation_split=0.1, shuffle=True, batch_size=batch_size,
Train on 62640 samples, validate on 6960 samples
Epoch 1/10
62640/62640 [=====] - 285s 5ms/step - loss: 0.0318 - acc: 0.9908 - val_loss: 0
.0078 - val_acc: 0.9990
Epoch 2/10
62640/62640 [=====] - 289s 5ms/step - loss: 0.0252 - acc: 0.9931 - val_loss: 0
.0201 - val_acc: 0.9953
Epoch 3/10
62640/62640 [=====] - 297s 5ms/step - loss: 0.0244 - acc: 0.9929 - val_loss: 0
.0229 - val_acc: 0.9941
Epoch 4/10
62640/62640 [=====] - 313s 5ms/step - loss: 0.0200 - acc: 0.9944 - val_loss: 0
.0271 - val_acc: 0.9922
Epoch 5/10
62640/62640 [=====] - 295s 5ms/step - loss: 0.0206 - acc: 0.9937 - val_loss: 0
.0122 - val_acc: 0.9954
Epoch 6/10
62640/62640 [=====] - 298s 5ms/step - loss: 0.0134 - acc: 0.9964 - val_loss: 0
.0036 - val_acc: 0.9994
Epoch 7/10
62640/62640 [=====] - 283s 5ms/step - loss: 0.0193 - acc: 0.9939 - val_loss: 0
.0135 - val_acc: 0.9955
Epoch 8/10
62640/62640 [=====] - 303s 5ms/step - loss: 0.0197 - acc: 0.9941 - val_loss: 0
.0154 - val_acc: 0.9955
Epoch 9/10
62640/62640 [=====] - 304s 5ms/step - loss: 0.0093 - acc: 0.9975 - val_loss: 0
.0066 - val_acc: 0.9989
Epoch 10/10
62640/62640 [=====] - 291s 5ms/step - loss: 0.0083 - acc: 0.9977 - val_loss: 9
.5236e-04 - val_acc: 0.9999
```

```
In [95]: y_pred = model.predict(x_test, batch_size=1024, verbose=1)
1040/1040 [=====] - 2s 2ms/step
```


```
In [96]: submission_df = pd.DataFrame(columns=['id'] + label_names)
submission_df['id'] = test_comments.index.values
submission_df[label_names] = y_pred
submission_df.to_csv("./lstm_submission.csv", index=False)
```

```
In [97]: from sklearn.metrics import classification_report
report = classification_report(np.argmax(y_test,axis =1), np.argmax(y_pred,axis =1))
```

```
In [98]: print(report)
```

	precision	recall	f1-score	support
0	0.61	0.34	0.43	140
1	0.78	0.93	0.85	500
2	0.94	0.86	0.90	400
micro avg	0.82	0.82	0.82	1040
macro avg	0.78	0.71	0.73	1040
weighted avg	0.82	0.82	0.81	1040


```
In [99]: import seaborn
from sklearn.metrics import confusion_matrix
confusion_matrix = confusion_matrix(np.argmax(y_test,axis =1),np.argmax(y_pred,axis =1))
matrix_proportions = np.zeros((3,3))
for i in range(0,3):
    matrix_proportions[i,:] = confusion_matrix[i,:]/float(confusion_matrix[i,:].sum())
names=['Hate','Offensive','Neither']
confusion_df = pd.DataFrame(matrix_proportions, index=names,columns=names)
plt.figure(figsize=(5, 5))
seaborn.heatmap(confusion_df,annot=True,annot_kws={"size": 12},cmap='YlGnBu',cbar=False, square=True,fmt
plt.ylabel(r'True categories',fontsize=14)
plt.xlabel(r'Predicted categories',fontsize=14)
plt.tick_params(labelsize=12)
```



In []:

In []:

In []:

In []: